

```
In [286]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
import statsmodels.regression.linear_model as lm
```

```
In [253]: data=pd.read_csv("Startups.csv")
data.head(10)
```

Out[253]:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96

```
In [254]: real_x=data.iloc[:,0:4].values
real_y=data.iloc[:,4].values
```

```
In [255]: le=LabelEncoder()
real_x[:,3]=le.fit_transform(real_x[:,3])
oneHE=OneHotEncoder()
real_x=oneHE.fit_transform(real_x).toarray()
```

```
In [256]: real_x=real_x[:,1:]
```

```
In [257]: training_x,test_x,training_y,test_y=train_test_split(real_x,real_y,test_size=
0.2,random_state=0)
```

```
In [258]: MLR=LinearRegression()
MLR.fit(training_x,training_y)
```

Out[258]: LinearRegression()

```
In [259]: pred_y=MLR.predict(test_x)
          pred_y
```

```
Out[259]: array([122629.67939821, 103103.18776136, 114542.21097071, 100892.89701574,
                116688.41484393, 117949.2268177 , 118090.69121311, 117930.63315589,
                112572.93581677, 116688.41484393])
```

```
In [260]: test_y
```

```
Out[260]: array([103282.38, 144259.4 , 146121.95,  77798.83, 191050.39, 105008.31,
                81229.06,  97483.56, 110352.25, 166187.94])
```

```
In [261]: MLR.coef_
```

```
Out[261]: array([-1.76704566e+04, -1.09433183e+04, -3.34903441e+04, -6.95039891e+03,
                -2.53951476e+03, -1.18483632e+04, -2.33160709e+04, -1.57955178e+04,
                -2.86591133e+04, -9.15369154e+03, -1.88504280e+04,  3.16023006e+03,
                 7.12202793e+03, -1.04765256e+03, -6.67528641e+03, -2.69957282e+03,
                -4.37127551e+02, -1.68249921e+03,  1.38302238e+04, -7.61040835e+03,
                 5.94126455e+03, -6.46567569e+03, -2.68097915e+03, -4.11547903e+03,
                -2.55228721e+03,  8.45721427e+03, -6.29746175e+03,  1.05841619e+04,
                -9.85050010e+02,  1.85907781e+04, -3.41489072e+02,  1.47381697e+04,
                 7.05705535e+03,  2.60212330e+02,  3.17759762e+03, -2.14620387e+03,
                 1.12091850e+03,  3.87297451e+03,  8.93643981e+03,  1.66114566e+04,
                 1.30213951e+04,  1.21203047e+04,  1.87323066e+04,  0.00000000e+00,
                 2.07572613e+04,  0.00000000e+00,  3.06221566e+04,  2.38772080e+04,
                -6.56834010e+03, -3.21507428e+04,  0.00000000e+00, -3.47010858e+01,
                 0.00000000e+00, -1.56363620e+04,  0.00000000e+00,  0.00000000e+00,
                -2.55551461e+03, -2.89035211e+03,  1.21203047e+04,  0.00000000e+00,
                -6.61710422e+03, -7.44374578e+02,  1.66114566e+04,  0.00000000e+00,
                 5.04561280e+03,  3.95998711e+03, -8.15684865e+03, -1.68536604e+04,
                -4.26220951e+04, -6.26599430e+03,  2.07572613e+04,  8.24158899e+03,
                 4.08295776e+03,  0.00000000e+00, -2.23804039e+04,  1.33944421e+03,
                 8.92002490e+03, -2.19604135e+04,  0.00000000e+00, -4.07625461e+02,
                -2.22302635e+04,  9.82179508e+03,  2.38772080e+04, -3.17573698e+03,
                -4.20129381e+03,  2.23997585e+03,  1.30213951e+04,  1.87323066e+04,
                 1.13225621e+04,  3.06221566e+04,  0.00000000e+00, -6.00757881e+03,
                 1.86914725e+04, -4.16606798e+03, -1.14384485e+04,  6.02063032e+03,
                -1.08256879e+03,  0.00000000e+00, -3.51355966e+04, -2.23804039e+04,
                -1.14384485e+04, -2.19604135e+04, -4.26220951e+04, -4.07625461e+02,
                 0.00000000e+00, -8.15684865e+03, -6.00757881e+03,  0.00000000e+00,
                 1.87323066e+04, -4.20129381e+03, -3.17573698e+03, -2.89035211e+03,
                -2.55551461e+03,  0.00000000e+00, -6.26599430e+03, -3.47010858e+01,
                 0.00000000e+00, -6.56834010e+03,  1.33944421e+03, -1.56363620e+04,
                -1.08256879e+03, -6.61710422e+03,  0.00000000e+00,  0.00000000e+00,
                 8.92002490e+03,  9.82179508e+03,  6.02063032e+03,  4.08295776e+03,
                 8.24158899e+03,  2.23997585e+03,  3.95998711e+03, -1.68536604e+04,
                 5.04561280e+03, -4.16606798e+03,  0.00000000e+00, -7.44374578e+02,
                 1.66114566e+04,  1.13225621e+04,  1.30213951e+04,  0.00000000e+00,
                 1.21203047e+04,  0.00000000e+00,  2.07572613e+04,  0.00000000e+00,
                 3.06221566e+04,  2.38772080e+04, -1.53834689e+04,  1.37935585e+03,
                 5.32114698e+03])
```

In [262]: `MLR.intercept_`

Out[262]: 115309.05898987639

In [263]: `real_x = np.append(arr=np.ones((50,1)).astype(int), values=real_x, axis=1)`

In [307]: `x_opt = real_x[:,[0,1,2,3,4,5]]`

In [308]: `# OLS=mls.ols( data,endog=real_y,exog=x_opt).fit()`

In [309]: `reg_OLS = sm.OLS(endog = real_y, exog = x_opt).fit()`

In [310]: `reg_OLS.summary()`

Out[310]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.203
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.113
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2.244
<b>Date:</b>	Fri, 26 Jun 2020	<b>Prob (F-statistic):</b>	0.0665
<b>Time:</b>	17:13:44	<b>Log-Likelihood:</b>	-594.98
<b>No. Observations:</b>	50	<b>AIC:</b>	1202.
<b>Df Residuals:</b>	44	<b>BIC:</b>	1213.
<b>Df Model:</b>	5		
<b>Covariance Type:</b>	nonrobust		

  

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.178e+05	5659.867	20.808	0.000	1.06e+05	1.29e+05
<b>x1</b>	-8.209e+04	3.84e+04	-2.139	0.038	-1.59e+05	-4730.369
<b>x2</b>	-5.284e+04	3.84e+04	-1.377	0.176	-1.3e+05	2.45e+04
<b>x3</b>	-6.828e+04	3.84e+04	-1.779	0.082	-1.46e+05	9086.971
<b>x4</b>	-4.801e+04	3.84e+04	-1.251	0.218	-1.25e+05	2.94e+04
<b>x5</b>	-3.654e+04	3.84e+04	-0.952	0.346	-1.14e+05	4.08e+04

  

<b>Omnibus:</b>	1.590	<b>Durbin-Watson:</b>	0.512
<b>Prob(Omnibus):</b>	0.452	<b>Jarque-Bera (JB):</b>	0.795
<b>Skew:</b>	-0.059	<b>Prob(JB):</b>	0.672
<b>Kurtosis:</b>	3.606	<b>Cond. No.</b>	7.47

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [311]: x_opt = real_x[:,[0,1,2,3,4,5]]
reg_OLS = sm.OLS(endog = real_y, exog = x_opt).fit()
reg_OLS.summary()
```

Out[311]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.203
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.113
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2.244
<b>Date:</b>	Fri, 26 Jun 2020	<b>Prob (F-statistic):</b>	0.0665
<b>Time:</b>	17:13:44	<b>Log-Likelihood:</b>	-594.98
<b>No. Observations:</b>	50	<b>AIC:</b>	1202.
<b>Df Residuals:</b>	44	<b>BIC:</b>	1213.
<b>Df Model:</b>	5		
<b>Covariance Type:</b>	nonrobust		

  

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.178e+05	5659.867	20.808	0.000	1.06e+05	1.29e+05
<b>x1</b>	-8.209e+04	3.84e+04	-2.139	0.038	-1.59e+05	-4730.369
<b>x2</b>	-5.284e+04	3.84e+04	-1.377	0.176	-1.3e+05	2.45e+04
<b>x3</b>	-6.828e+04	3.84e+04	-1.779	0.082	-1.46e+05	9086.971
<b>x4</b>	-4.801e+04	3.84e+04	-1.251	0.218	-1.25e+05	2.94e+04
<b>x5</b>	-3.654e+04	3.84e+04	-0.952	0.346	-1.14e+05	4.08e+04

  

<b>Omnibus:</b>	1.590	<b>Durbin-Watson:</b>	0.512
<b>Prob(Omnibus):</b>	0.452	<b>Jarque-Bera (JB):</b>	0.795
<b>Skew:</b>	-0.059	<b>Prob(JB):</b>	0.672
<b>Kurtosis:</b>	3.606	<b>Cond. No.</b>	7.47

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [316]:

```
x_opt = real_x[:,[0,1,2,3,4]]
reg_OLS = sm.OLS(endog = real_y, exog = x_opt).fit()
reg_OLS.summary()
```

Out[316]:

OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.187
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.115
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2.584
<b>Date:</b>	Fri, 26 Jun 2020	<b>Prob (F-statistic):</b>	0.0496
<b>Time:</b>	17:14:21	<b>Log-Likelihood:</b>	-595.48
<b>No. Observations:</b>	50	<b>AIC:</b>	1201.
<b>Df Residuals:</b>	45	<b>BIC:</b>	1211.
<b>Df Model:</b>	4		
<b>Covariance Type:</b>	nonrobust		

  

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.17e+05	5592.160	20.917	0.000	1.06e+05	1.28e+05
<b>x1</b>	-8.13e+04	3.83e+04	-2.121	0.039	-1.59e+05	-4083.603
<b>x2</b>	-5.205e+04	3.83e+04	-1.358	0.181	-1.29e+05	2.52e+04
<b>x3</b>	-6.748e+04	3.83e+04	-1.760	0.085	-1.45e+05	9733.737
<b>x4</b>	-4.721e+04	3.83e+04	-1.232	0.225	-1.24e+05	3e+04

  

<b>Omnibus:</b>	1.208	<b>Durbin-Watson:</b>	0.460
<b>Prob(Omnibus):</b>	0.547	<b>Jarque-Bera (JB):</b>	0.481
<b>Skew:</b>	-0.012	<b>Prob(JB):</b>	0.786
<b>Kurtosis:</b>	3.480	<b>Cond. No.</b>	7.38

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [317]: x_opt = real_x[:,[0,1,2,3]]
reg_OLS = sm.OLS(endog = real_y, exog = x_opt).fit()
reg_OLS.summary()
```

Out[317]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.159
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.105
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2.908
<b>Date:</b>	Fri, 26 Jun 2020	<b>Prob (F-statistic):</b>	0.0445
<b>Time:</b>	17:14:34	<b>Log-Likelihood:</b>	-596.31
<b>No. Observations:</b>	50	<b>AIC:</b>	1201.
<b>Df Residuals:</b>	46	<b>BIC:</b>	1208.
<b>Df Model:</b>	3		
<b>Covariance Type:</b>	nonrobust		

  

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.16e+05	5563.332	20.845	0.000	1.05e+05	1.27e+05
<b>x1</b>	-8.03e+04	3.85e+04	-2.083	0.043	-1.58e+05	-2710.722
<b>x2</b>	-5.104e+04	3.85e+04	-1.324	0.192	-1.29e+05	2.65e+04
<b>x3</b>	-6.648e+04	3.85e+04	-1.725	0.091	-1.44e+05	1.11e+04

  

<b>Omnibus:</b>	0.787	<b>Durbin-Watson:</b>	0.365
<b>Prob(Omnibus):</b>	0.675	<b>Jarque-Bera (JB):</b>	0.207
<b>Skew:</b>	0.028	<b>Prob(JB):</b>	0.902
<b>Kurtosis:</b>	3.310	<b>Cond. No.</b>	7.30

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [318]: x_opt = real_x[:,[0,1,3]]
reg_OLS = sm.OLS(endog = real_y, exog = x_opt).fit()
reg_OLS.summary()
```

Out[318]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.127
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.090
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	3.430
<b>Date:</b>	Fri, 26 Jun 2020	<b>Prob (F-statistic):</b>	0.0407
<b>Time:</b>	17:14:55	<b>Log-Likelihood:</b>	-597.25
<b>No. Observations:</b>	50	<b>AIC:</b>	1200.
<b>Df Residuals:</b>	47	<b>BIC:</b>	1206.
<b>Df Model:</b>	2		
<b>Covariance Type:</b>	nonrobust		

  

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.149e+05	5549.041	20.707	0.000	1.04e+05	1.26e+05
<b>x1</b>	-7.923e+04	3.88e+04	-2.040	0.047	-1.57e+05	-1089.548
<b>x2</b>	-6.541e+04	3.88e+04	-1.684	0.099	-1.44e+05	1.27e+04

  

<b>Omnibus:</b>	0.458	<b>Durbin-Watson:</b>	0.343
<b>Prob(Omnibus):</b>	0.796	<b>Jarque-Bera (JB):</b>	0.074
<b>Skew:</b>	0.061	<b>Prob(JB):</b>	0.964
<b>Kurtosis:</b>	3.144	<b>Cond. No.</b>	7.22

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



```
In [319]: x_opt = real_x[:,[0,1]]
reg_OLS = sm.OLS(endog = real_y, exog = x_opt).fit()
reg_OLS.summary()
```

Out[319]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.075
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.055
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	3.875
<b>Date:</b>	Fri, 26 Jun 2020	<b>Prob (F-statistic):</b>	0.0548
<b>Time:</b>	17:15:18	<b>Log-Likelihood:</b>	-598.71
<b>No. Observations:</b>	50	<b>AIC:</b>	1201.
<b>Df Residuals:</b>	48	<b>BIC:</b>	1205.
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

  

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.136e+05	5596.183	20.294	0.000	1.02e+05	1.25e+05
<b>x1</b>	-7.79e+04	3.96e+04	-1.969	0.055	-1.57e+05	1665.637

  

<b>Omnibus:</b>	0.172	<b>Durbin-Watson:</b>	0.216
<b>Prob(Omnibus):</b>	0.918	<b>Jarque-Bera (JB):</b>	0.033
<b>Skew:</b>	0.061	<b>Prob(JB):</b>	0.984
<b>Kurtosis:</b>	2.966	<b>Cond. No.</b>	7.15

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [320]:

```
x_opt = real_x[:,[0]]
reg_OLS = sm.OLS(endog = real_y, exog = x_opt).fit()
reg_OLS.summary()
```

Out[320]:

OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.000
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.000
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	nan
<b>Date:</b>	Fri, 26 Jun 2020	<b>Prob (F-statistic):</b>	nan
<b>Time:</b>	17:15:35	<b>Log-Likelihood:</b>	-600.65
<b>No. Observations:</b>	50	<b>AIC:</b>	1203.
<b>Df Residuals:</b>	49	<b>BIC:</b>	1205.
<b>Df Model:</b>	0		
<b>Covariance Type:</b>	nonrobust		

  

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	1.12e+05	5700.155	19.651	0.000	1.01e+05	1.23e+05

  

<b>Omnibus:</b>	0.018	<b>Durbin-Watson:</b>	0.020
<b>Prob(Omnibus):</b>	0.991	<b>Jarque-Bera (JB):</b>	0.068
<b>Skew:</b>	0.023	<b>Prob(JB):</b>	0.966
<b>Kurtosis:</b>	2.825	<b>Cond. No.</b>	1.00

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.